Ziniu Li

ziniuli@link.cuhk.edu.cn
http://www.liziniu.org
https://scholar.google.com/citations?user=8oUnKQQAAAAJ&hl=en
https://www.zhihu.com/people/mo-fei-10-41
https://x.com/ZiniuLi

Research

My research focuses on the **algorithm design** and **theoretical analysis** of **machine learning** models, particularly in (large-scale) **reinforcement learning**. Currently, I am primarily working on **large language models** to pursue the development of an agent capable of universal intelligence. My work has been recognized with several honors, including **Best-Paper Runner-Up** (at NeurIPS 2024 FITML Workshop), **Oral Presentation** (ICLR2024 Tiny Paper Track, UAI 2023, NeurIPS 2021 EcoRL Workshop), and **Spotlight Presentation** (NeurIPS 2023). Below, I highlight key contributions from my previous research:

- In the field of large language models, my work spans several key areas: data selection (NeurIPS 2023 Spotlight), diversity-preserving supervised fine-tuning (ICLR 2025, NeurIPS 2024 FITML Workshop Best Paper Runner-up), generalization of RLHF (ICLR2024 Tiny Paper Oral), computationally efficient RLHF (ICML 2024), and hallucination mitigation (ICLR2 2025).
- In the field of imitation learning and reinforcement learning, I am interested in the theory of sample complexity (NeurIPS 2020, TPAMI 2021, UAI 2023 Oral), efficient exploration (ICLR 2022, NeurIPS 2021 EcoRL Workshop Oral, DAI 2020), as well as applications in robotics (ICLR 2024 Blog) and signal processing (TSP 2024).
- I also work on optimization-centric topics with other researchers, including understanding Adam in training Transformers (NeurIPS 2024), memory-efficient optimizers (ICLR 2025), zero-order optimization (IJCAI 2020), and prompt-tuning (EMNLP 2024 Finding).

Education

08/2020 - Present — **Ph.D.** at The Chinese University of Hong Kong, Shenzhen
Advisor: Zhi-Quan (Tom) Luo

08/2015 – 06/2019 — **B.E.** at Xi'an Jiaotong University
Advisor: Zhiyuan Liu

Employment History

10/2021 – 08/2022 — **Algorithm Engineer (Intern)** at Tencent, Shenzhen
Advisor: Peilin Zhao

06/2020 – 09/2020 — **Algorithm Engineer (Intern)** at Cardinal Operations, Shanghai

07/2019 – 06/2020 — **Research Assistant** at Nanjing University, Nanjing
Advisor: Yang Yu

12/2018 – 02/2019 — **Algorithm Engineer (Intern)** at Tianrang Technology Inc., Hangzhou
Advisor: Zhenhui (Jessie) Li

Selected Publications

* indicates equal contribution.

1. **Ziniu Li**, Congliang Chen, Tian Xu, Zeyu Qin, Jiancong Xiao, Zhi-Quan Luo, Ruoyu Sun. *Preserving Diversity in Supervised Fine-tuning of Large Language Models.* In Proceedings of the 13th International Conference on Learning Representations (ICLR), 2025.

- This paper is selected as an **best paper runner-up** at NeurIPS FITML Workshop, 2024.

2. **Ziniu Li**, Tian Xu, Yushun Zhang, Zhihang Lin, Yang Yu, Ruoyu Sun, Zhi-Quan Luo. *ReMax: A Simple, Effective, and Efficient Reinforcement Learning Method for Aligning Large Language Models.* In Proceedings of the 41st Conference on International Conference on Machine Learning (ICML), 2024.

3. **Ziniu Li***, Tian Xu*, Yang Yu. *When is RL better than DPO in RLHF? A Representation and Optimization Perspective.* In Proceedings of the 12th International Conference on Learning Representations (ICLR) (Tiny Paper Track), 2024.

   - This paper is selected as an **oral** presentation at ICLR, 2024.

4. **Ziniu Li***, Tian Xu*, Zeyu Qin, Yang Yu, Zhi-Quan Luo. *Imitation Learning from Imperfection: Theoretical Justifications and Algorithms.* Advances in on Neural Information Processing System 37 (NeurIPS), 2023.

   - This paper is selected as an **spotlight** presentation (acceptance rate $< 5\%$) at NeurIPS, 2023.

5. Tian Xu*, **Ziniu Li***, Yang Yu, Zhi-Quan Luo. *Provably Efficient Adversarial Imitation Learning with Unknown Transitions.* In Proceedings of the 39th Conference on Uncertainty in Artificial Intelligence (UAI), 2023.

   - This paper is selected as an **oral** presentation (acceptance rate $< 3\%$) at UAI, 2023.

6. **Ziniu Li**, Yingru Li, Yushun Zhang, Tong Zhang, Zhi-Quan Luo. *HyperDQN: A Randomized Exploration Method for Deep Reinforcement Learning.* In Proceedings of the 10th International Conference on Learning Representations (ICLR), 2022.

   - This paper is selected as an **oral** presentation in Workshop on Ecological Theory of Reinforcement Learning at NeurIPS, 2021.

7. Tian Xu, **Ziniu Li**, Yang Yu. *Error Bounds of Imitating Policies and Environments for Reinforcement Learning.* IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI), 2021.

Full Publications

Journal Articles

[Fan+24] 1 Youlin Fan, Bo Jiu, Wenqiang Pu, **Ziniu Li**, Kang Li, and Hongwei Liu. "Sensing Jamming Strategy From Limited Observations: An Imitation Learning Perspective". In: *IEEE Transactions on Signal Processing* (2024).

[XLY21] 2 Tian Xu, **Ziniu Li**, and Yang Yu. "Error bounds of imitating policies and environments for reinforcement learning". In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 44.10 (2021), pp. 6968–6980.

Conference Proceedings

[Jia+25] 1 Chengxing Jia, **Ziniu Li**, Pengyuan Wang, Yi-Chen Li, Zhenyu Hou, Yuxiao Dong, and Yang Yu. "Controlling Large Language Model with Latent Actions". In: *International Conference on Machine Learning*. 2025.

[Li+25] 2 **Ziniu Li**, Congliang Chen, Tian Xu, Zeyu Qin, Jiancong Xiao, Zhi-Quan Luo, and Ruoyu Sun. "Preserving Diversity in Supervised Fine-tuning of Large Langguage Models". In: *International Conference on Learning Representations*. 2025.

[Yan+25] 3 Tianyun Yang, **Ziniu Li**, Juan Cao, and Chang Xu. "Understanding and Mitigating Hallucination in Large Vision-Language Models via Modular Attribution and Intervention". In: *International Conference on Learning Representations*. 2025.

[Zha+25] 4 Yushun Zhang, Congliang Chen, **Ziniu Li**, Tian Ding, Chenwei Wu, et al. "Adam-mini: Use fewer learning rates to gain more". In: *International Conference on Learning Representations*. 2025.

[LXY24] 5 **Ziniu Li**, Tian Xu, and Yang Yu. "When is RL better than DPO in RLHF? A Representation and Optimization Perspective". In: *The Second Tiny Papers Track of International Conference on Learning Representations*. 2024.

[Yan+24] 6 Tianyun Yang, **Ziniu Li**, Juan Cao, and Chang Xu. "Pruning for Robust Concept Erasing in Diffusion Models". In: *Neurips Safe Generative AI Workshop 2024*. 2024.

[Zha+24a] 7 Heshen Zhan, Congliang Chen, Tian Ding, **Ziniu Li**, and Ruoyu Sun. "Unlocking Black-Box Prompt Tuning Efficiency via Zeroth-Order Optimization". In: *Findings of the Association for Computational Linguistics: EMNLP 2024*. 2024, pp. 14825–14838.

[Zha+24b] 8 Yushun Zhang, Congliang Chen, Tian Ding, **Ziniu Li**, Ruoyu Sun, and Zhi-Quan Luo. "Why transformers need adam: A hessian perspective". In: *Advances in Neural Information Processing Systems 37*. 2024.

[Li+23b] 9 **Ziniu Li**, Tian Xu, Zeyu Qin, Yang Yu, and Zhi-Quan Luo. "Imitation learning from imperfection: Theoretical justifications and algorithms". In: *Advances in Neural Information Processing Systems 36*. 2023.

[Li+23c] 10 **Ziniu Li**, Tian Xu, Yushun Zhang, Zhihang Lin, Yang Yu, Ruoyu Sun, and Zhi-Quan Luo. "Remax: A simple, effective, and efficient reinforcement learning method for aligning large language models". In: *International Conference on Machine Learning*. 2023.

[Xu+23] 11 Tian Xu, **Ziniu Li**, Yang Yu, and Zhi-Quan Luo. "Provably Efficient Adversarial Imitation Learning with Unknown Transitions". In: *Proceedings of the 39th Conference on Uncertainty in Artificial Intelligence*. 2023, pp. 2367–2378.

[Li+22] 12 **Ziniu Li**, Tian Xu, Yang Yu, and Zhi-Quan Luo. "Rethinking valuedice-does it really improve performance". In: *ICLR Blog Track*. 2022.

[Li+21] 13 **Ziniu Li**, Yingru Li, Yushun Zhang, Tong Zhang, and Zhi-Quan Luo. "HyperDQN: A randomized exploration method for deep reinforcement learning". In: *International Conference on Learning Representations*. 2021.

[LC20] 14 **Ziniu Li** and Xiong-Hui Chen. "Efficient Exploration by Novelty-Pursuit". In: *Proceedings of the 2nd International Conference on Distributed Artificial Intelligence*. 2020, pp. 85–102.

[LLQ20] 15 Fei-Yu Liu, **Zi-Niu Li**, and Chao Qian. "Self-Guided Evolution Strategies with Historical Estimated Gradients." In: *International Joint Conferences on Artificial Intelligence*. 2020, pp. 1474–1480.

[XLY20] 16 Tian Xu, **Ziniu Li**, and Yang Yu. "Error Bounds of Imitating Policies and Environments". In: *Advances in Neural Information Processing Systems 33*. 2020, pp. 15737–15749.

arXiv Preprints

[Tan+25a] 1 Zhengyang Tang, **Ziniu Li**, Zhenyang Xiao, Tian Ding, Ruoyu Sun, et al. *Enabling Scalable Oversight via Self-Evolving Critic*. In: *arXiv preprint arXiv:2501.05727* (2025).

[Tan+25b] 2 Zhengyang Tang, **Ziniu Li**, Zhenyang Xiao, Tian Ding, Ruoyu Sun, et al. *RealCritic: Towards Effectiveness-Driven Evaluation of Language Model Critiques*. In: *arXiv preprint arXiv:2501.14492* (2025).

[Jia+24] 3 Chengxing Jia, Pengyuan Wang, **Ziniu Li**, Yi-Chen Li, Zhilong Zhang, Nan Tang, and Yang Yu. *BWArea Model: Learning World Model, Inverse Dynamics, and Policy for Controllable Language Generation*. In: *arXiv preprint arXiv:2405.17039* (2024).

[Xia+24] 4 Jiancong Xiao, **Ziniu Li**, Xingyu Xie, Emily Getzen, Cong Fang, Qi Long, and Weijie J Su. *On the Algorithmic Bias of Aligning Large Language Models with RLHF: Preference Collapse and Matching Regularization*. In: *arXiv preprint arXiv:2405.16455* (2024).

[Li+23a] 5 **Ziniu Li**, Ke Xu, Liu Liu, Lanqing Li, Deheng Ye, and Peilin Zhao. *Deploying Offline Reinforcement Learning with Human Feedback*. In: *arXiv preprint arXiv:2303.07046* (2023).

[LXY22] 6 **Ziniu Li**, Tian Xu, and Yang Yu. *A Note on Target Q-learning For Solving Finite MDPs with A Generative Oracle*. In: *arXiv preprint arXiv:2203.11489* (2022).

[Xu+22] 7 Tian Xu, **Ziniu Li**, Yang Yu, and Zhi-Quan Luo. *Understanding Adversarial Imitation Learning in Small Sample Regime: A Stage-coupled Analysis*. In: *arXiv preprint arXiv:2208.01899* (2022).

Service

- Organizers: Reinforcement Learning Weekly Seminar, Deep Learning Seminar at CUHKSZ.
- T-PAMI, NeurIPS'2024, ICML'2024, ICLR'2024, NeurIPS'2023, ICML'2023, NeurIPS'2022 (Top Reviewer), ICML'2022 (Outstanding Reviewer), ICLR'2022 (Highlighted Reviewer), DAI'2020.
- Technical Program Committee Members: RL4RealLife@NeurIPS2022, DAI'2022.

Awards

2024 Best paper runner-up award at NeurIPS 2024 Workshop on Fine-Tuning in Modern Machine Learning: Principles and Scalability

2021 Best paper presentation award at the Doctoral and Postdoctoral Forum of Shenzhen Research Institute of Big Data

2017 Honor student of Xi'an Jiaotong University

Second prize of National Electrical Math Modeling Competition

2016 Second prize of Chinese Mathematics Competition, Shaanxi Province